

Objective

Main objective of the current project is to build a image classifier using Convoluted neural networks & through traditional fully connected DNNs. To understand what makes CNNs special. This model can be able to separate out multiple classes of images and may be useful in image gallery sorting or searching of images etc.

Importing libraries

```
In [58]: from tensorflow import keras
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import os
from tensorflow.keras import layers
from tensorflow.keras import Sequential
import cv2
import numpy as np
import warnings
from sklearn.metrics import classification_report
import gc
from numba import cuda
warnings.filterwarnings("ignore")
%matplotlib inline
```

Loading data

```
In [2]: train_folder = os.path.abspath('./data/seg_train/seg_train/')
test_folder = os.path.abspath('./data/seg_test/seg_test/')


```

Helper functions

```
In [3]: def read_image(folder, file_name):
    return cv2.imread(os.path.join(folder, file_name), cv2.IMREAD_UNCHANGED)

def join_directories(folders):
    current_path = None
    for folder in folders:
        if current_path:
            current_path = os.path.join(current_path, folder)
        else:
            current_path = folder
        continue
    return current_path
```

Description

- The data chosen for this project is a publicly available data set provided by Intel corporation and available at <https://www.kaggle.com/puneet6060/intel-image-classification>
- This Data contains around 25k images of size 150x150 distributed under 6 categories.
 - Buildings
 - Forest
 - Glacier
 - Mountain
 - Sea
 - Street

```
In [4]: folder_sizes = dict()
for folder in os.listdir(train_folder):
    folder_sizes[folder] = len(os.listdir(os.path.join(train_folder, folder)))

print("Label distribution is as follows")
display(folder_sizes)
```

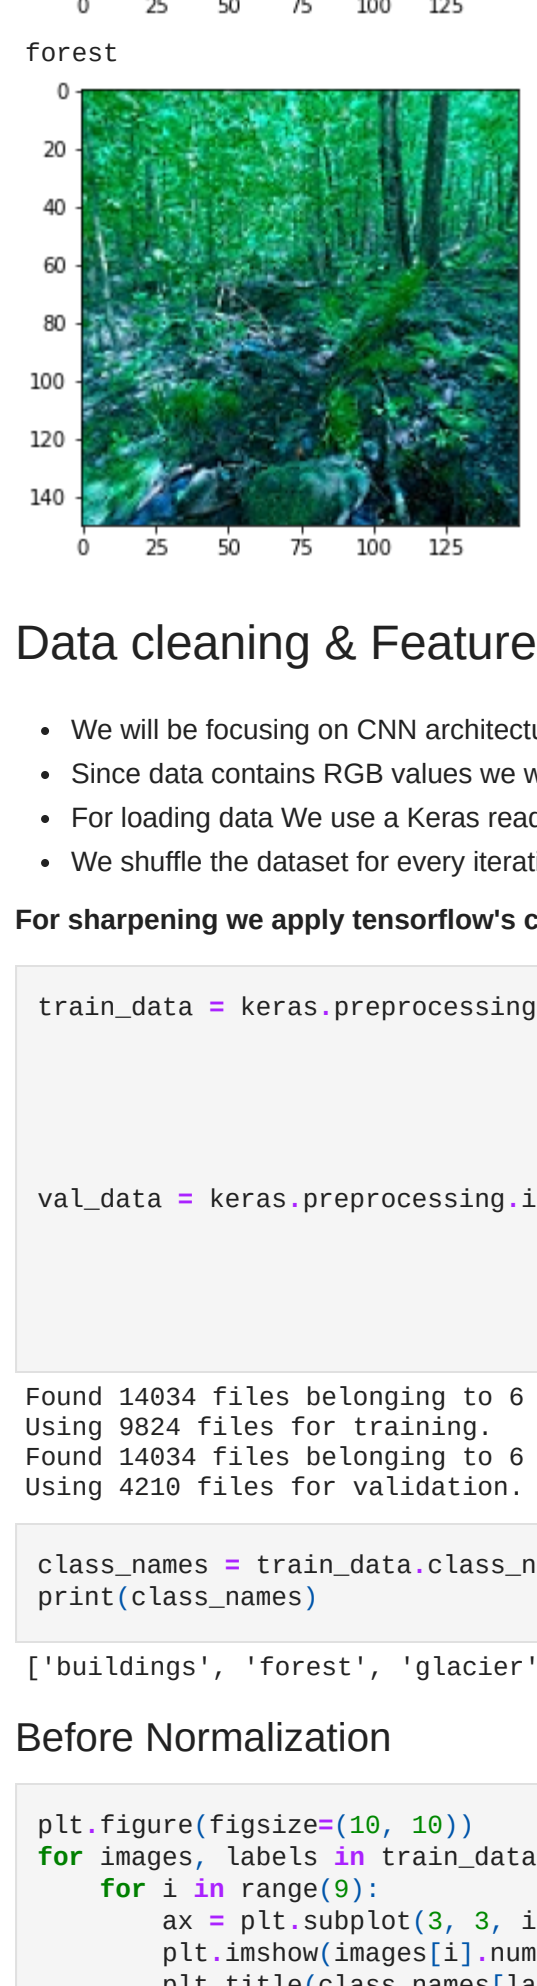
Label distribution is as follows

```
{'buildings': 2191,
 'street': 2282,
 'mountain': 2512,
 'glacier': 2484,
 'sea': 2274,
 'forest': 2271}
```

```
In [5]: label_files = dict()
for folder in os.listdir(train_folder):
    label_files[folder] = [file for file in os.listdir(os.path.join(train_folder, folder))]
```

Sample Images

```
In [6]: for label in label_files:
    plt.imshow(read_image(join_directories([train_folder, label]), label_files[label][0]))
    plt.show()
```



Data cleaning & Feature engineering

- We will be focusing on CNN architectures to derive relevant features for us.
- Since data contains RGB values we will be Normalizing them.
- For loading data We use a Keras reading library which can convert our data to batches of Tensors
- We shuffle the dataset for every iteration

For sharpening we apply tensorflow's contrast adjustment module

```
In [16]: train_data = keras.preprocessing.image_dataset_from_directory(train_folder,
                                                                    labels='infer',
                                                                    image_size = (150, 150), shuffle=True,
                                                                    validation_split=0.3,
                                                                    subset='training', seed=0, batch_size=16)

val_data = keras.preprocessing.image_dataset_from_directory(train_folder,
                                                            labels='infer',
                                                            image_size = (150, 150), shuffle=True,
                                                            validation_split=0.3,
                                                            subset='validation', seed=0, batch_size=16)

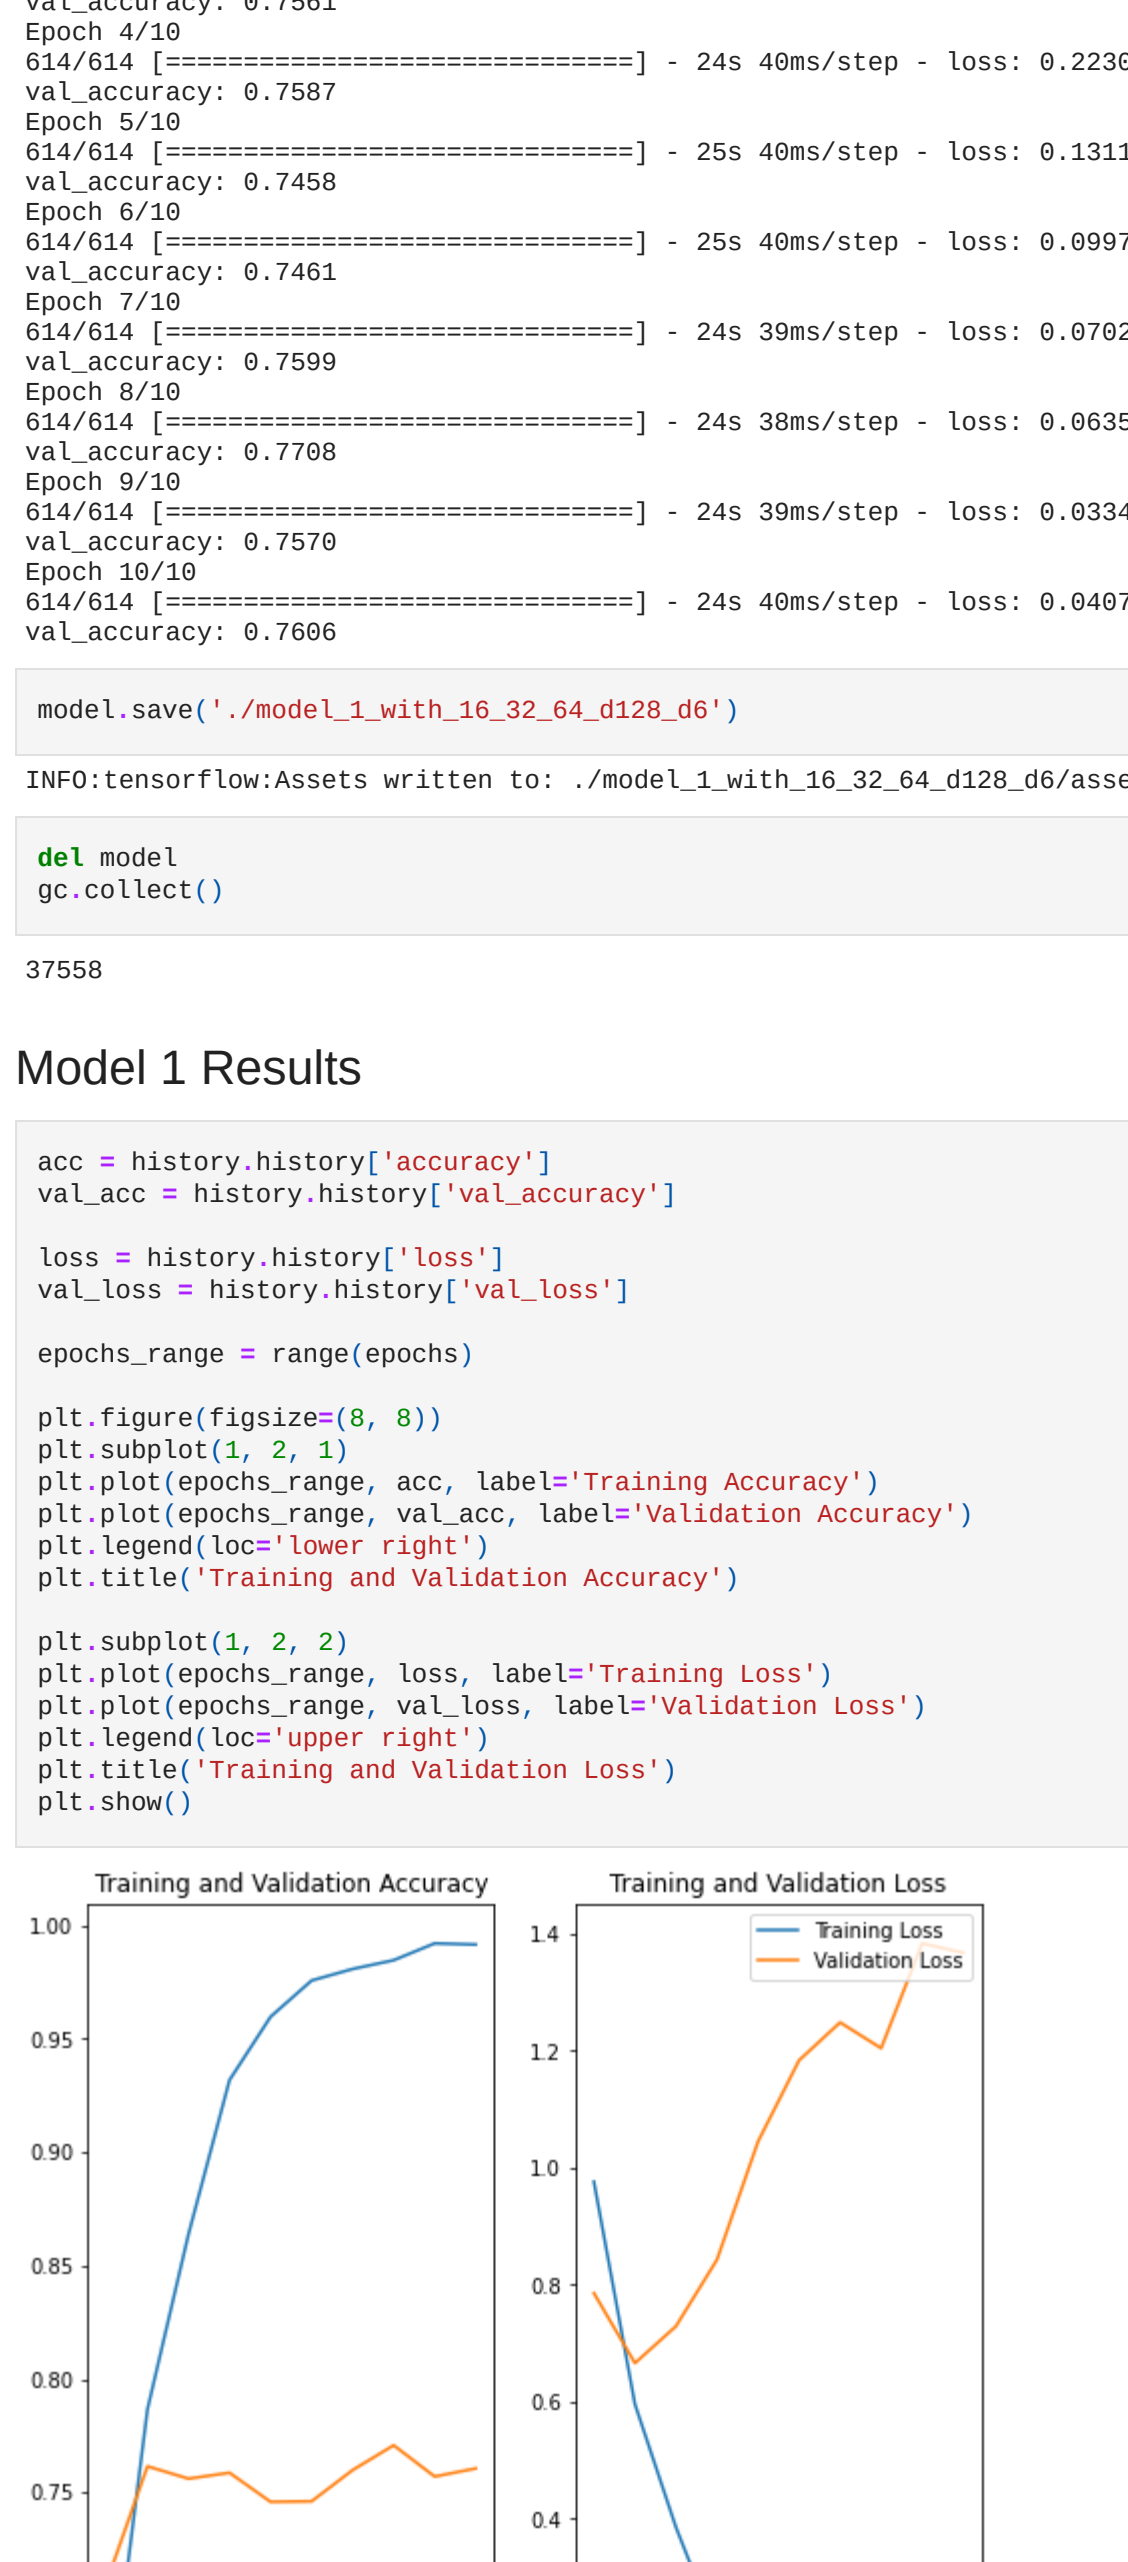
Found 14934 files belonging to 6 classes.
Using 9824 files for training.
Found 14934 files belonging to 6 classes.
Using 4218 files for validation.
```

```
In [7]: class_names = train_data.class_names
print(class_names)

['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

Before Normalization

```
In [55]: plt.figure(figsize=(10, 10))
Epoch 5/10
for i in range(9):
    ax = plt.subplot(3, 3, 1 + i)
    plt.imshow(images[i].numpy().astype("uint8"))
    plt.title(class_names[labels[i]])
    plt.axis("off")
```



Training

```
In [6]: num_classes = 6
epochs=10
```

Model - 1

- In this model we will be using Convolution 2D layers with Max pooling. With a single Dense layer before output layer

```
In [60]: model = Sequential([
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(150, 150, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    # layers.Conv2D(64, 3, padding='same', activation='relu'),
    # layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

In [61]: model.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
In [62]: model.summary()

Model: "sequential_3"
Layer (type) Output Shape Param #
=====
rescaling_4 (Rescaling) (None, 150, 150, 3) 0
conv2d_9 (Conv2D) (None, 150, 150, 16) 448
max_pooling2d_9 (MaxPooling2D) (None, 75, 75, 16) 0
conv2d_10 (Conv2D) (None, 75, 75, 32) 4640
max_pooling2d_10 (MaxPooling2D) (None, 37, 37, 32) 0
flatten_3 (Flatten) (None, 43808) 0
dense_6 (Dense) (None, 128) 5687552
dense_7 (Dense) (None, 6) 774
=====
Total params: 5,613,414
Trainable params: 5,613,414
Non-trainable params: 0
```

```
In [63]: history = model.fit(
    train_data,
    validation_data=val_data,
    epochs=epochs
)

Epoch 1/10
614/614 [=====] - 131s 212ms/step - loss: 1.3344 - accuracy: 0.5427 - val_loss: 0.7855
val_accuracy: 0.7189
Epoch 2/10
614/614 [=====] - 24s 39ms/step - loss: 0.6529 - accuracy: 0.8460 - val_loss: 0.6660
val_accuracy: 0.7615
Epoch 3/10
614/614 [=====] - 24s 39ms/step - loss: 0.4226 - accuracy: 0.9055 - val_loss: 0.7290
val_accuracy: 0.7561
Epoch 4/10
614/614 [=====] - 24s 40ms/step - loss: 0.2260 - accuracy: 0.9264 - val_loss: 0.8428
val_accuracy: 0.7587
Epoch 5/10
614/614 [=====] - 25s 40ms/step - loss: 0.1311 - accuracy: 0.9564 - val_loss: 1.0445
val_accuracy: 0.7458
Epoch 6/10
614/614 [=====] - 25s 40ms/step - loss: 0.0997 - accuracy: 0.9689 - val_loss: 1.1834
val_accuracy: 0.7461
Epoch 7/10
614/614 [=====] - 24s 39ms/step - loss: 0.0702 - accuracy: 0.9804 - val_loss: 1.2481
val_accuracy: 0.7599
Epoch 8/10
614/614 [=====] - 24s 38ms/step - loss: 0.0635 - accuracy: 0.9784 - val_loss: 1.2848
val_accuracy: 0.7788
Epoch 9/10
614/614 [=====] - 24s 39ms/step - loss: 0.0334 - accuracy: 0.9927 - val_loss: 1.3832
val_accuracy: 0.7245
Epoch 10/10
614/614 [=====] - 24s 40ms/step - loss: 0.0487 - accuracy: 0.9984 - val_loss: 1.3672
val_accuracy: 0.7686
```

```
In [65]: model.save('./model_1_with_16_32_64_d128_d6')
```

INFO:tensorflow:Assets written to: ./model_1_with_16_32_64_d128_d6/assets

```
In [67]: del model
gc.collect()
```

Out[67]: 37558

Model 1 Results


```
In [64]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [69]: device = cuda.get_current_device()
device.reset()
```

Model 2

- In this model we will by the same network but this time to reduce the overfitting we can try doing a random dropout

```
In [8]: model_2 = Sequential([
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(150, 150, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    # layers.Conv2D(64, 3, padding='same', activation='relu'),
    # layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

In [10]: model_2.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
In [11]: model_2.summary()

Model: "sequential_4"
Layer (type) Output Shape Param #
=====
rescaling_1 (Rescaling) (None, 150, 150, 3) 0
conv2d_2 (Conv2D) (None, 150, 150, 16) 448
max_pooling2d_2 (MaxPooling2D) (None, 75, 75, 16) 0
conv2d_3 (Conv2D) (None, 75, 75, 32) 4640
max_pooling2d_3 (MaxPooling2D) (None, 37, 37, 32) 0
dropout_1 (Dropout) (None, 37, 37, 32) 0
flatten_1 (Flatten) (None, 43808) 0
dense_1 (Dense) (None, 128) 5687552
dense_2 (Dense) (None, 6) 774
=====
Total params: 5,613,414
Trainable params: 5,613,414
Non-trainable params: 0
```

```
In [13]: history_2 = model_2.fit(
    train_data,
    validation_data=val_data,
    epochs=epochs
)

Epoch 1/10
614/614 [=====] - 31s 43ms/step - loss: 1.3312 - accuracy: 0.5199 - val_loss: 0.7631
val_accuracy: 0.7276
Epoch 2/10
614/614 [=====] - 24s 39ms/step - loss: 0.6887 - accuracy: 0.8240 - val_loss: 0.6325
val_accuracy: 0.6922
Epoch 3/10
614/614 [=====] - 24s 39ms/step - loss: 0.4749 - accuracy: 0.8787 - val_loss: 0.6932
val_accuracy: 0.7553
Epoch 4/10
614/614 [=====] - 24s 40ms/step - loss: 0.2813 - accuracy: 0.9044 - val_loss: 0.9063
val_accuracy: 0.7572
Epoch 5/10
614/614 [=====] - 24s 40ms/step - loss: 0.1742 - accuracy: 0.9435 - val_loss: 1.0556
val_accuracy: 0.7705
Epoch 6/10
614/614 [=====] - 24s 39ms/step - loss: 0.1215 - accuracy: 0.9625 - val_loss: 1.1314
val_accuracy: 0.7722
Epoch 7/10
614/614 [=====] - 24s 40ms/step - loss: 0.0917 - accuracy: 0.9714 - val_loss: 1.1769
val_accuracy: 0.7720
Epoch 8/10
614/614 [=====] - 24s 39ms/step - loss: 0.0566 - accuracy: 0.9831 - val_loss: 1.3447
val_accuracy: 0.7553
Epoch 9/10
614/614 [=====] - 25s 40ms/step - loss: 0.0681 - accuracy: 0.9806 - val_loss: 1.5231
val_accuracy: 0.7245
Epoch 10/10
614/614 [=====] - 24s 39ms/step - loss: 0.0587 - accuracy: 0.9815 - val_loss: 1.4486
val_accuracy: 0.7722
```

```
In [14]: model_2.save('./model_2_with_16_32_64_drop_d128_d6')
```

INFO:tensorflow:Assets written to: ./model_2_with_16_32_64_drop_d128_d6/assets

```
In [15]: del model_2
gc.collect()
```

Out[15]: 6814

Model 2 results

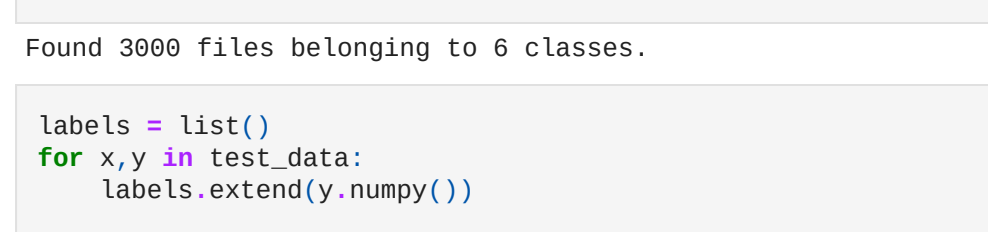
```
In [16]: acc = history_2.history['accuracy']
val_acc = history_2.history['val_accuracy']

loss = history_2.history['loss']
val_loss = history_2.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [33]: device = cuda.get_current_device()
device.reset()
```

Model 3

- As Overfitting is not getting resolved we can try increasing the dataset by Data augmentation techniques such as random rotation, translation etc.

```
In [7]: datagen = ImageDataGenerator(
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    validation_split=0.3,
)
```

```
In [8]: train_data = datagen.flow_from_directory(
    train_folder,
    target_size = (150, 150), shuffle=True,
    subset='training', seed=0, batch_size=16,
    class_mode='binary')

val_data = datagen.flow_from_directory(
    train_folder,
    target_size = (150, 150), shuffle=True,
    subset='validation', seed=0, batch_size=16,
    class_mode='binary')

Found 9826 images belonging to 6 classes.
Found 4208 images belonging to 6 classes.
```

```
In [9]: model_3 = Sequential([
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(150, 150, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    # layers.Conv2D(64, 3, padding='same', activation='relu'),
    # layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

In [10]: model_3.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
In [11]: model_3.summary()

Model: "sequential_5"
Layer (type) Output Shape Param #
=====
rescaling (Rescaling) (None, 150, 150, 3) 0
conv2d (Conv2D) (None, 150, 150, 16) 448
max_pooling2d (MaxPooling2D) (None, 75, 75, 16) 0
conv2d_1 (Conv2D) (None, 75, 75, 32) 4640
max_pooling2d_1 (MaxPooling2D) (None, 37, 37, 32) 0
dropout (Dropout) (None, 37, 37, 32) 0
flatten (Flatten) (None, 43808) 0
dense (Dense) (None, 128) 5687552
dense_1 (Dense) (None, 6) 774
=====
Total params: 5,613,414
Trainable params: 5,613,414
Non-trainable params: 0
```

```
In [12]: history_3 = model_3.fit(
    train_data,
    validation_data=val_data,
    epochs=epochs
)

Epoch 1/10
615/615 [=====] - 106s 166ms/step - loss: 1.3872 - accuracy: 0.4813 - val_loss: 0.9166
val_accuracy: 0.6419
Epoch 2/10
615/615 [=====] - 98s 169ms/step - loss: 0.8223 - accuracy: 0.7019 - val_loss: 0.7232
val_accuracy: 0.7395
Epoch 3/10
615/615 [=====] - 100s 163ms/step - loss: 0.7225 - accuracy: 0.7370 - val_loss: 0.7170
val_accuracy: 0.7293
Epoch 4/10
615/615 [=====] - 101s 164ms/step - loss: 0.6381 - accuracy: 0.7703 - val_loss: 0.5943
val_accuracy: 0.7866
Epoch 5/10
615/615 [=====] - 101s 164ms/step - loss: 0.5743 - accuracy: 0.7937 - val_loss: 0.5638
val_accuracy: 0.7997
Epoch 6/10
615/615 [=====] - 100s 163ms/step - loss: 0.5325 - accuracy: 0.8060 - val_loss: 0.6303
val_accuracy: 0.7797
Epoch 7/10
615/615 [=====] - 100s 163ms/step - loss: 0.4954 - accuracy: 0.8223 - val_loss: 0.5485
val_accuracy: 0.7745
Epoch 8/10
615/615 [=====] - 100s 162ms/step - loss: 0.4677 - accuracy: 0.8279 - val_loss: 0.5345
val_accuracy: 0.8082
Epoch 9/10
615/615 [=====] - 101s 164ms/step - loss: 0.4311 - accuracy: 0.8419 - val_loss: 0.5385
val_accuracy: 0.8032
Epoch 10/10
615/615 [=====] - 100s 163ms/step - loss: 0.4311 - accuracy: 0.8415 - val_loss: 0.5283
val_accuracy: 0.8158
```

```
In [13]: model_3.save('./model_2_with_horflip_shear2_zoom_2_16_32_64_drop_d128_d6')
INFO:tensorflow:Assets written to: ./model_2_with_horflip_shear2_zoom_2_16_32_64_drop_d128_d6/assets
```

```
In [14]: del model_3
gc.collect()
```

Out[14]: 6795

Model 3 results

```
In [15]: acc = history_3.history['accuracy']
val_acc = history_3.history['val_accuracy']

loss = history_3.history['loss']
val_loss = history_3.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [17]: test_data = keras.preprocessing.image_dataset_from_directory(test_folder,
                                                                    labels='infer',
                                                                    image_size = (150, 150), shuffle=False,
                                                                    seed=0, batch_size=16)

Found 9900 files belonging to 6 classes.
```

```
In [49]: labels = list()
for x,y in test_data:
    labels.extend(y.numpy())
```

```
In [54]: model_1_4 = tf.keras.models.load_model('./model_1_with_16_32_64_d128_d6/')
del model_1
gc.collect()
```

Out[54]: 10707

```
In [55]: model_2 = tf.keras.models.load_model('./model_2_with_16_32_64_drop_d128_d6/')
model_2_predictions = model_2.predict_classes(test_data)
del model_2
gc.collect()
```

Out[55]: 11713

```
In [56]: model_3 = tf.keras.models.load_model('./model_2_with_horflip_shear2_zoom_2_16_32_64_drop_d128_d6/')
model_3_predictions = model_3.predict_classes(test_data)
del model_3
gc.collect()
```

Out[56]: 11713

```
In [66]: print("Model 1 Classification report")
print(classification_report(labels,model_1_predictions))
```

Model 1 Classification report

	precision	recall	f1-score	support
0	0.76	0.68	0.72	437
1	0.93	0.93	0.93	474
2	0.79	0.59	0.68	553
3	0.63	0.78	0.70	525
4	0.79	0.71	0.75	519
5	0.78	0.80	0.79	501
accuracy			0.75	3800
macro avg	0.76	0.75	0.75	3800
weighted avg	0.76	0.75	0.75	3800

```
In [67]: print("Model 2 Classification report")
print(classification_report(labels,model_2_predictions))
```

Model 2 Classification report

	precision	recall	f1-score	support
0	0.75	0.72	0.74	437
1	0.93	0.93	0.93	474
2	0.79	0.59	0.73	553
3	0.69	0.78	0.73	525
4	0.74	0.71	0.73	519
5	0.78	0.80	0.79	501
accuracy			0.77	3800
macro avg	0.77	0.77	0.77	3800
weighted avg	0.77	0.77	0.77	3800

```
In [68]: print("Model 3 Classification report")
print(classification_report(labels,model_3_predictions))
```

Model 3 Classification report

	precision	recall	f1-score	support
0	0.87	0.72	0.79	437
1	0.87	0.98	0.92	474
2	0.79	0.82	0.75	525
3	0.70	0.76	0.73	525
4	0.81	0.76	0.78	519
5	0.94	0.85	0.89	491
accuracy			0.80	3800
macro avg	0.81	0.81	0.81	3800
weighted avg	0.81	0.80	0.80	3800

Summary & Key findings

- Based on the training curves obtained from the models, it is evident that a lot of overfitting is happening very early in the model itself.
- Adding dropout increased max acquired Val accuracy slightly 0.77 with drop out & 0.75 without
- Augmenting the data improved the test accuracy further to 0.80 and validation loss was decreasing(More scope for better results).
- Based on these observations & the classification reports shown above we should go with the Model 3 - with Data augmentation.

Further steps

- As Augmenting the data is improving the Validation set accuracy, adding more data will definitely result in better performance of the model.
- Augmenting data also improved the training a lot, as the training curves indicate both training accuracy & validation accuracy are improving. We can train for more epochs which will definitely result in better accuracy.
- As overfitting is happening quite early in the model, we can decrease the complexity of the model. Reducing no. of convolution layers will help.
- Hyper parameter tuning such as changing the optimizer, changing batch size might help better generalize. Due to GPU restrictions of My laptop I am not able to test with higher batch size than 16. And any model above 10 Lakh trainable parameters is becoming unwieldy for the GPU.

In []: